

CS 6120: Comparing Models for Multilabel Classification of PubMed Article Abstracts

Nithish Bhat and Ashish Thomas and Jamie Tjia

Northeastern University

Boston, MA, United States

{bhat.nithi, thomas.ash, tjia.j}@northeastern.edu

Abstract

In this paper, we discuss our project on multi-label classification of PubMed articles into root labels using NLP. We approach the biomedical classification task with BERT (BioBERT) and compare it against three baseline models: RNN, CNN, and Logistic Regression.

Code used in this project available [here](#).

1 Introduction

The growth of biomedical literature has created a need for automated methods to organize and classify research articles effectively. Manually labeling articles with relevant Medical Subject Headings (MeSH) is labor-intensive and time-consuming, making automated multi-label classification a helpful tool for researchers and clinicians.

This project aims to use a transformer-based language model, specifically BERT, to perform multi-label classification on biomedical research articles. Instead of traditional single-label classification, we use multi-label classification, as articles are typically associated with multiple MeSH tags.

Our goals are: (1) to train and evaluate a BERT model for accurately predicting MeSH root labels and (2) to compare its performance against more conventional models, including logistic regression, CNNs, and RNNs. By doing so, we hope to demonstrate the advantages of transformer-based models in capturing contextual information and label correlations in biomedical texts, ultimately providing a tool for organizing and analyzing the PubMed corpus.

In our classification task, we encounter challenges such as fixing class imbalance, dealing with specialized language, and optimizing training time / compute. Different evaluation metrics and source data resizing became part of the solution to some of these problems. Through this report, we detail our data processing, model architectures, experi-

mental setup, and evaluation results, along with an analysis of each approach.

2 Background and Related Work

The rapid rise in biomedical publications makes it hard for scientists to find useful data. To help, the National Library of Medicine (NLM) tags PubMed entries using Medical Subject Headings (MeSH), a structured list grouping medical ideas in levels. Yet, tagging by hand takes much effort, money, and time—making expansion tough (Koutsomitropoulos and Andriopoulos, 2020). For this reason, automatic Multi-Label Text Classification (MLTC) has become vital in bio-NLP tasks. In contrast to typical single-label sorting, MLTC links several tags to one document, since most health papers cover various subjects at once—for example Disorders, Substances, or Treatment Methods (Tsoumakas and Katakis, 2009).

Before Transformers became common, models for text analysis often used CNNs instead of RNNs. While both handled sequences differently, they formed the base for early deep learning in this area. CNNs started in image processing but later got used for text tasks after Kim’s work (Kim, 2014), using small sliding filters to spot short meaningful word groups. Although good at finding fixed-length phrases, they usually fail to link distant parts of longer texts because their focus area is narrow. RNNs plus Bi-LSTMs: recurrent systems like Bidirectional Long Short-Term Memory (Bi-LSTM) emerged to handle text step by step. Because they analyze data from left to right and from right to left, they grasp sequence patterns along with word positioning better than fixed bag-of-words methods (Hochreiter and Schmidhuber, 1997). Still, such architectures run into processing delays thanks to their linear flow; also, gradients tend to weaken across extended chains.

Devlin et al. (Devlin et al., 2019) introduced BERT along with this new structure, changing how

NLP works. Instead of using RNNs, these models rely on self-attention to handle input at once—enabling better understanding of context and links between labels.

Although standard BERT works well on common tasks, it struggles with biomedical texts because of differences in terminology - like drug names or clinical terms. Instead of using general data, Lee et al. (Lee et al., 2019) trained BioBERT on specialized sources such as PubMed and PMC, keeping the original structure. This version performs better in spotting entities, finding links between them, also answering medical questions. Here, we use biobert-base-cased-v1.2 to identify specific patterns tied to MeSH labels.

3 Data

The baseline and BERT models were all trained and tested on the dataset [PubMed MultiLabel Text Classification Dataset MeSH](#). This dataset contains information on 50,000 articles from PubMed, including each article’s title, abstract, and Medical Subject Heading (MeSH) labels.

MeSH labels are arranged into a nested hierarchy with 16 root labels (see Table 1). Two labels—Humanities [K] and Publication Characteristics [V]—have been excluded due to their infrequency in the available data. The 14 included labels each appear in at least ten percent of the sample articles (see Figure 1).

Each article has up to 13 root labels, with a median root label count of six (see Figure 2). The article abstracts are an average of 222 tokens long (see Figure 3).

4 Methods

We adjusted the BioBERT model (biobert-base-cased-v1.2) to assign medical summaries to up to 14 MeSH groups that may overlap. Texts were processed using BioBERT’s cased WordPiece tokenizer, which inserts special [CLS] and [SEP] tokens and pads or truncates sequences to a fixed length of 128 tokens. Instead of stacking multiple layers, a single linear module was connected to the model’s final output, transforming hidden states into 14 separate prediction values. This meant the final hidden state of just the [CLS] token was passed through a single linear layer that produced 14 independent logits. Since articles might fit more than one label at once, outputs used individual sigmoid functions per node, avoiding softmax which

assumes exclusive classes. This setup allowed parallel probability estimates across topics. Training relied on BCEWithLogitsLoss—a unified method combining sigmoids and binary loss—for better computational reliability during optimization. The entire BioBERT encoder and classification head were fine-tuned jointly using the AdamW optimizer. At prediction time, a fixed cutoff of 0.5 was used; categories surpassing the cutoff were linked to the text.

The performance of the fine-tuned BERT model was measured against three baseline models which were trained as follows.

The logistic regression model serves as a basic neural starting point while performing effectively like a "Neural Bag-of-Words" system. Although it uses an embedding layer set to 256 dimensions by default, this component converts word IDs into compact vector forms. Rather than tracking sequence position, it computes the average of these embedded values throughout the full text entry. As a result, one consolidated mean vector emerges, capturing the overall meaning of the abstract. This average form feeds straight into a linear layer, then uses BCEWithLogitsLoss for training. The approach estimates each label’s likelihood separately, yet maintains fast response times.

The CNN approach detects local word sequences and meaningful phrases in text, no matter where they appear. Instead of focusing on exact positions, it processes embedded inputs with length 256 using a 1D convolutional layer with a window size of 50 and moves forward by steps of 20 to capture distinct signals from the abstract. Following this step, global max pooling selects the strongest activation per filter over the full input span. Then, these selected outputs use BCEWithLogitsLoss for training so that multiple labels can be predicted at once.

The RNN model handles word order by applying a Bi-LSTM setup. First, an embedding step converts split text into 256 element vectors. Then, these representations enter a Bi-LSTM block containing 128 memory cells. Instead of one direction only, processing occurs left-to-right and right-to-left to grasp full-sentence meaning. Afterward, outputs from both flows merge before entering a dense layer. That last component computes unnormalized scores directly. During inference, logits become probabilities using a sigmoid activation. With a set cutoff, multi-label outputs turn into binary results.




Figure 1: Frequency of MeSH Root Labels




Figure 2: Number of Labels per Abstract

5 Experiments

We used micro F1, macro F1, weighted-average F1, exact match ratio, and Hamming loss to evaluate the performance of each model. Micro F1 calculates the F1 score using precision and recall

over all of the labels at once, while the macro F1 calculates the F1 for each label and returns the average across the labels. Similarly to macro F1, weighted-average F1 calculates the F1 for each label and returns the weighted average. The exact match ratio measures how many predictions exactly




Figure 3: Distribution of Abstract Lengths

matched every label, and Hamming loss measures how many labels were predicted incorrectly across all labels. Higher F1 scores and exact match ratios indicate better performance, and lower Hamming loss indicates better performance.

The learning rate for each model was determined with separate tests for each model on the dev dataset. All of the models were trained over 25 epochs, using an initial learning rate of 0.001 for the baselines and 0.0001 for the BERT model with Adam optimization, on both reduced-size and full-size versions of the dataset. At the end of each epoch, the evaluation metrics were calculated for on dev dataset to determine the ideal number of epochs for tests.

5.1 Results

Looking at the results, the BioBERT model leads baseline models in performance on the classification task. Micro and Macro F1 are higher for BioBERT compared to the baseline models, which shows that even with class imbalance BioBERT beats out the baseline. The amount by which BioBERT beats all three baselines is more for the smaller dataset, with performance at around 0.7-0.8, and this shows the difference between the BERT fine-tuning vs. training from start in the baseline models. This difference across reduced

and full data continues across to exact match ratio and Hamming loss: BioBERT has lower Hamming loss and higher exact match ratio than baseline.

Although it is the most straightforward model, linear regression the best performance of the baseline models. Linear regression had worse performance than RNN with regard to F1 scores the first several epochs of training, the model using linear regression had better performance with regard to all metrics past 7 epochs of training.

RNN had the best initial performance of the baseline models, but peaked in performance about 5 epochs into training, allowing linear regression to surpass RNN.

The CNN model had by far the worst performance, with lower F1 scores, lower exact match ratios, and higher Hamming loss for the majority of training. CNN only performed better than logistic regression for the first one or two epochs, and never performed better than RNN for any evaluation metric.

6 Conclusions

From this project, it is clear BioBERT performs better than the baseline models. This illustrates the importance of properly pretraining embeddings, particularly when the subject matter requires expert knowledge to understand, as with biomedical




Figure 5: Performance Metrics by Epoch

research. Even with large amounts of training data and several epochs of training, the baseline models approach the performance of BioBERT but do not surpass it.

The next steps for this work would be to assess the performance of BERT-based models for the more complicated task of predicting the full nested MeSH labels of PubMed articles. In particular, evaluating the performance of a more naive BERT model against both hierarchical multilabel classification models—which model the nested hierarchy of the labels when making predictions—and a model that uses BERT embeddings alongside hierarchical multilabel classification model (Zhang et al., 2025). Another way of expanding this project would be to compare BioBERT, a word2vec model

trained on biomedical text, and similar models trained on less specific datasets, in order to assess the importance of pretraining embeddings on text which uses the highly specialized vocabulary used in biomedical research papers.

References

- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Sepp Hochreiter and Jürgen Schmidhuber. 1997.

Long short-term memory. *Neural Comput.*, 9(8):1735–1780.

Yoon Kim. 2014. Convolutional neural networks for sentence classification. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1746–1751, Doha, Qatar. Association for Computational Linguistics.

Dimitrios A. Koutsomitropoulos and Andreas D. Andriopoulos. 2020. Automated mesh indexing of biomedical literature using contextualized word representations. In *Artificial Intelligence Applications and Innovations*, pages 343–354, Cham. Springer International Publishing.

Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. 2019. BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4):1234–1240. [_eprint: https://academic.oup.com/bioinformatics/article-pdf/36/4/1234/48983216/bioinformatics_36_4_1234.pdf](https://academic.oup.com/bioinformatics/article-pdf/36/4/1234/48983216/bioinformatics_36_4_1234.pdf).

Grigorios Tsoumakas and Ioannis Katakis. 2009. Multi-label classification: An overview. *International Journal of Data Warehousing and Mining*, 3:1–13.

Wei Zhang, Yun Jiang, Yun Fang, and Shuai Pan. 2025. Hierarchical contrastive learning for multi-label text classification. *Scientific Reports*, 15.

Label	Name
A	Anatomy
B	Organisms
C	Diseases
D	Chemicals and Drugs
E	Analytical, Diagnostic and Therapeutic Techniques, and Equipment
F	Psychiatry and Psychology
G	Phenomena and Processes
H	Disciplines and Occupations
I	Anthropology, Education, Sociology, and Social Phenomena
J	Technology, Industry, and Agriculture
K	Humanities
L	Information Science
M	Named Groups
N	Health Care
V	Publication Characteristics
Z	Geographicals

Table 1: Full Label Names